

COT 830 - Homework 2

Problem 1 - Distributed Mutual Exclusion

- (1) In Rickart-Agrawala's algorithm, critical sections are accessed according to the increasing order of timestamps. We prove this using contradiction.
- Assume the contrary to above, i.e. suppose process p_1 issues a request to enter critical section at time t_1 & process p_2 does so at time t_2 with $t_1 < t_2$ & p_2 enters CS first.
- From algorithm, p_2 enters CS $\Rightarrow p_2$ received all reply msgs $\Rightarrow p_2^P$ receives reply msg from p_1 .
- p_1 reply to p_2 's CS request $\rightarrow (p_1 \text{ not requesting CS}) \vee (t_2 < t_1)$
- Since p_1 requested CS, $(p_1 \text{ not req. CS})$ is false
- from Modus ponens, if assumption is true & assumption $\rightarrow (t_2 < t_1)$ then $t_2 < t_1$. But this is a contradiction! as $t_1 < t_2$ was the assumption (part of it)
- ∴ $\neg(\text{Assumption})$ is true always & hence CS are accessed in incr. order of timestamps

(2) Modified Ricart Agrawala algorithm

- Each process P_i keep some additional state
 - $\text{pred}_i \leftarrow$ process id of immediate predecessor (i.e. executes CS just before that / priority of req. is just higher than P_i)
 - $\text{next}_i \leftarrow$ process id of imm. successor
- Requesting CS as P_0
 - (a) IA periodically Broadcast (REQ) message to all processes.
 - (b) When P_j recv REQ msg.
 - If $t_{\text{pred}} < t_j$ (t_j is timestamp of REQ) $\& (t_i < t_j)$
 $t_{\text{pred}} = t_j$; $\text{pred}_j = i$
 - Else If $(P_j \text{ req. CS}) \& (t_i > t_j) \& (t_i < t_{\text{next}} - \epsilon_j)$
 $t_{\text{next}} = t_i$; $\text{next}_i = j$
 - Executing CS for P_i
 - If $[(P_i \text{ recv. REPLY from pred}) \& (\text{pred} \neq \text{null})] \text{ OR}$
 $[(P_i \text{ recv. REPLY from all proc.}) \& (\text{pred} = \text{null})]$
then execute (CS_P)
 - Releasing CS for P_i
 - Send REPLY to next_j , reset $t_{\text{next}} = \text{INT-MIN}$

→ Correctness :

- (i) 'pred_i' has process id of process which has requested the last amongst all processes that requested before P_i. Thus pred_i gives process with priority just higher than i's req.
- (ii) 'next_i' has process id of process which has requested the CS when after i's request. A. It is the process with highest priority / least timestamp amongst all such processes. Since t_i < t_{next}, i's priority just precedes next_i & so it will be waiting for reply to request from i only.
- (iii) Initially, if no one makes requests & Only i does, it will need a reply from all processes, ~~even if~~ ^{even if} process ~~unless~~ unless someone requests CS, their req. time stamp can be higher, in which case it becomes successor, or lower, then it can become predecessor & i will need a reply.
- (iv) Since this only modifies no. of message, the earlier proof of access to CS in increasing order of timestamps, combined with (iii) also works here. Thus fairness is ensured.
- (v) Since fairness enforces an order of access of CS & since 2 requests can't have same priority (tie-break using proc. id) 2 proc can't enter CS simultaneously, ensuring correctness.
- (vi) Message complexity is $O(N-1) + 1 \approx O(N)$

- (3) In Lamport's algorithm, REPLY msg serves to flush all request messages in the channel with smaller timestamps. If P_i has recvd. REPLY from P_j, it has also recvd. all requests made by j at smaller timestamp & can decide which of these requests, or its own have higher priority & execute CS first.

- (a) For each process in rectangular grid P_{ij} , denoting the i^{th} row, j^{th} column (of numCol columns & numRow rows)
- $$R_{ij} = \{P_{ik}\}_{k=1}^{\text{numCol}} \cup \{P_{kj}\}_{k=1}^{\text{numRows}}, \text{ clearly } P_{ij} \in R_{ij}$$
- $R_{ij} \cap R_{i'j'} \neq \emptyset, = \{P_{ij'}, P_{i'j}\}$
 - $\forall i, j, i', j' \quad R_{ij} \neq R_{i'j'} \text{ since } (i, j) \text{ uniquely determines a proc}$
 - $|R_{ij}| = \text{numRows} + \text{numCols} - 1 = k \neq i, j$
 - $\forall P_{ij}, \text{ such } R^{ij} = \{R_{i'j'} \mid P_{ij} \in R_{i'j'}\}, |R^{ij}| = \text{numRow} + \text{numCol} - 1$
i.e. P_{ij} appears in exactly k req. sets
 - Given this new R_i set satisfies properties, we can run Mackawa's algorithm for mutual exclusion as it is, resulting in
 - Achieving mutual exclusion - proof by contradiction
 if P_{ij} & P_{kl} are in CS \rightarrow both received reply from R_{ij} & R_{kl} respectively $\rightarrow r \in R_{ij} \cap R_{kl}$ replied to requests made & since $R_{ij} \cap R_{kl} \neq \emptyset$, a process replied to 2 processes \rightarrow A contradiction
 - $3K$ messages - K REQ, K REPLY & K RELEASE with $k = \text{numRows} + \text{numCols} - 1$
 - Problem of Deadlocks - consider P_{ij} & P_{kl} req. CS & P_{il} replies to P_{ij} & P_{kj} replies to P_{kl} ($\{P_{il}, P_{kj}\} = R_{ij} \cap R_{kl}$)
Clearly there is a deadlock
 - Deadlocks can be handled like by Mackawa's extensions by using additional msg. types FAILED, INQUIRE & YIELD
so msg. complexity $\sim 6 \cdot 3K - 6K$

Problem 2 - Global Snapshots

- (1) We can modify Chandy-Lamport's snapshotting algo. as follows
- Marker sending Rule → same as original
 - Marker Receiving Rule → for process j : P_j
On receiving marker along channel C
If j hasn't recorded its state then
 - Record the state of $C = \emptyset \cup S_j$
 - Follow 'Marker sending Rule'
 - Else
let older snapshot = O_{Sj}
 - Make a new snapshot by combining $O_{Sj} + \text{current state}$ where current state is achieved after receiving all the messages, since last snapshot (O_{Sj})
 - Record current state of $C = \emptyset \cup S$
 - Termination - ~~$\forall i \neq j$~~ if i has seen marker along each incoming channel, it can exchange local snapshot w/ neighbours to aggregate or send it to its parent, which will do that.

- (2) Lai-Yang's algorithm, white \leq red. Now consider events s, t
 $\Rightarrow s \rightarrow t$. Possible cases are
- s, t on same process, so s can be white/red & t also white/red but if t is white, s can't be red $\Rightarrow s.\text{color} \leq t.\text{color}$
 - $s = \text{send}$ & $t = \text{matching Recv}$ If t is red $s.\text{color} \leq t.\text{color}$
If t is white & s is white, no issue $s.\text{color} = t.\text{color}$
~~If t is white & s is red isn't possible as t 's process turns red on receipt of red msg(s) & so $t.\text{color} = \text{red}$~~

$\rightarrow \exists u \ni s \rightarrow u \rightarrow t$. From above 2 cases

$$s.\text{color} \leq u.\text{color} \quad \& \quad u.\text{color} \leq t.\text{color} \Rightarrow s.\text{color} \leq t.\text{color}$$

$\rightarrow \exists u_1, u_2, \dots, u_n \ni s \rightarrow u_1 \rightarrow u_2 \rightarrow \dots \rightarrow u_n \rightarrow t$

By an inductive argument & transitivity of \leq , $s.\text{color} \leq t.\text{color}$

Problem 3 - Termination

(1) Assuming a controlling agent monitors computation & sends only control messages to processes, we can adapt the weight throwing algorithm to become ~~fault-tolerant~~^{message-loss} as follows -

\rightarrow Rule 1 - same as given in Kshemkalyani (controller throwing ~~its~~ a part of its weight to a process via basic msg). - P.T.O.

\rightarrow Rule 2 - on Receipt of message basic(DW) at proc j from i
If i has already sent basic(DW) i.e. j received same timestamp on a basic message from i w/ same weight
then ~~send~~ send ACK(timestamp(basic(DW))) to i

Else

$$w_j := w_j + DW$$

If P_j is idle, P_j becomes active

send ACK(timestamp(basic(DW))) to i

\rightarrow Rule 3 - Process i becomes idle from active state by sending control message ($DW = w_i$) to controlling agent

Wait(ACK(C(DW)));

If (timeout) resnd ~~the~~ C(DW) to controller & Wait();

If Recv(ACK(C(DW))) from controller

$w_i = 0$, change state to idle

\rightarrow Rule 4 Term - on Recv. C(DW), controller $w_c := w_c + DW$

If $w=1$, termination

Date 92 10 18

To expand Rule 1 →

Send $b(DW)$ to a process P_i

Wait $(ACK(b(DW)))$

If (timeout), resend $b(DW)$ to P_i & wait

If $Recv(ACK(b(DW)))$ then

$$w_c := w_c - DW$$

(2) We can modify Huang's weight throwing algorithm to use counter

→ For each process & controller, we keep ^{one} counter

- ^{Decrement} counter when (basic) message is sent to process

- ^{Increment} counter when (control) message is received from process

- If $\# i \min(counter(P_i), controller) = -k$, then the algorithm

is essentially the weight throwing algorithm, with $w = k$

→ we can argue termination in a similar manner. if at a process

P_i , $counter = 0$, then either process has not started yet

(since on receipt of messages, it becomes active & counter > 0)

or has moved its work to its children (+ve counter reduced to zero after sends) or has terminated & given it to controller

If the controller process' counter = 0, then algo may have not started yet, or has finished (equivalent to $w = 1$)